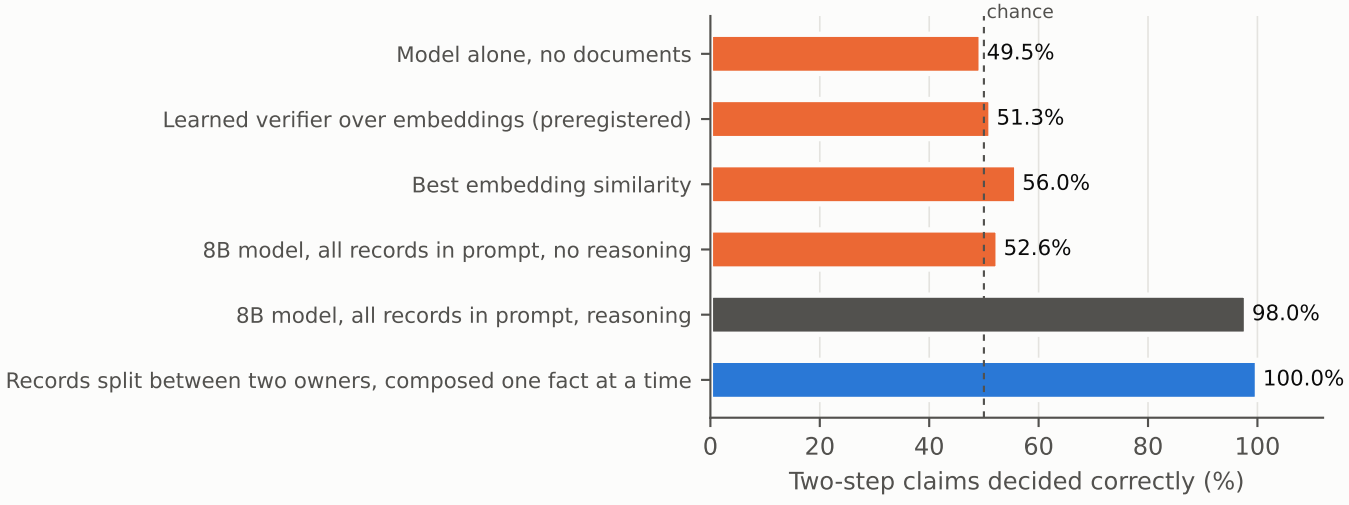
chance
Model alone, no documents
49.5%
Learned verifier over embeddings (preregistered)
51.3%
Best embedding similarity
56.0%
8B model, all records in prompt, no reasoning
52.6%
8B model, all records in prompt, reasoning
98.0%
Records split between two owners, composed one fact at a time
100.0%
0
20
40
60
80
100
Two-step claims decided correctly (%)
End-to-end neural
Needs every record in one place
Across owners, clean records